

Mnemosyne — Video Script (4 minutes)

[0:00–0:20] Hook — The Problem

[Screen: Generic LLM chat interface, blank text box]

Most LLM tools give you a blank text box and leave you to figure out what to ask. You type “build an app” and get a generic plan. Skip one important question and the whole thing falls apart.

[Screen: Fade to Mnemosyne logo]

Mnemosyne takes the opposite approach. It structures the conversation so that by the time you reach the planning stage, every assumption has been questioned and every risk has been surfaced.

[0:20–0:50] Setup — Auth & Configuration

[Screen: Login page, clean monospace UI]

Sign up with just a username and password. No email required.

[Screen: Settings modal, provider preset dropdown]

Open System Config, pick your provider from the dropdown — OpenAI, Anthropic, Groq, DeepSeek, and more — paste your API key, and you’re ready. The app works with any OpenAI-compatible endpoint.

[0:50–1:40] Stage 1–2: Intake and Clarify

[Screen: New chat, user types “I want to build a fitness tracking app for runners”]

Type your idea and hit send. The Coordinator agent reads it, generates a summary, and automatically advances to the Clarification stage.

You’ll see a loading indicator — “THINKING...” — then “LOADING CLARIFY...” as the system chains stages together without pause.

[Screen: Dynamic clarification form appears with context-specific questions]

The Inquisitor generates questions specific to YOUR idea. For a fitness app, it asks about target platforms, whether you need social features, and how you plan to monetize. Every question includes a short note explaining why it matters.

[Screen: User selects radio buttons, adjusts slider, checks multiselect options]

Answer using the form — radio buttons, sliders, text fields, yes/no toggles, and multi-select checkboxes. The input types are chosen by the AI based on what makes sense for each question.

[Screen: Submit, follow-up questions appear]

After submitting, the AI asks follow-up questions to dig deeper. When you're done, click "Skip and Proceed" or let it auto-advance after two rounds.

[1:40–2:20] Stage 3: Research

[Screen: Risk matrix appears]

The Analyst generates a risk matrix — real risks specific to your idea, each rated by impact and likelihood, with concrete mitigations.

[Screen: User clicks "Accept and Proceed"]

You decide: accept these risks and continue, or pivot and go back to clarifications. This is a human decision. The AI informs it but doesn't make it.

[2:20–2:50] Stage 4: Compile

[Screen: Compiled Brief card appears]

The Compiler gathers everything — your idea, your clarification answers, the risk assessment — and synthesizes it into a single structured brief.

[Screen: User reviews and clicks "Accept Brief and Finalize"]

Review it, then send it to the final stage.

[2:50–3:40] Stage 5: Finalizer

[Screen: Kanban execution board with 4 phases, column navigation arrows]

The Architect generates a full execution plan: Foundation, Build, Launch, and Grow. Each phase has 3-6 actionable tasks with details and deliverables.

[Screen: User checks off tasks one by one, progress bar fills]

Check off tasks as you complete them. The progress bar tracks your completion.

[Screen: Last task checked, card fills black, celebration animation]

When every task is done, the card fills black and a “PLAN COMPLETE!” celebration appears.

[Screen: User clicks pencil icon to uncheck a task]

Made a mistake? The edit button lets you reset and re-check.

[3:40–4:00] Wrap-up

[Screen: Sidebar showing multiple sessions, rename icon]

All your sessions are listed in the sidebar. Rename them inline, delete old ones, start new ones.

[Screen: Final shot of the full app]

Mnemosyne turns a rough idea into a structured plan by making sure nothing gets skipped. The conversation has a shape, the output has rigor, and the final plan is something you can actually execute.

[End]